

Open the dataset `ames_housing_no_missing.csv` with the following command:

```
import pandas as pd

ames_housing = pd.read_csv(
    "../datasets/ames_housing_no_missing.csv",
    na_filter=False,  # required for pandas>2.0
)

target_name = "SalePrice"

data = ames_housing.drop(columns=target_name)

target = ames_housing[target_name]
```

`ames_housing` is a pandas dataframe. The column "SalePrice" contains the target variable.

To simplify this exercise, we will only used the numerical features defined below:

```
numerical_features = [
    "LotFrontage", "LotArea", "MasVnrArea", "BsmtFinSF1", "BsmtFinSF2",
    "BsmtUnfSF", "TotalBsmtSF", "1stFlrSF", "2ndFlrSF", "LowQualFinSF",
    "GrLivArea", "BedroomAbvGr", "KitchenAbvGr", "TotRmsAbvGrd",
    "Fireplaces",
    "GarageCars", "GarageArea", "WoodDeckSF", "OpenPorchSF",
    "EnclosedPorch",
    "3SsnPorch", "ScreenPorch", "PoolArea", "MiscVal",
]

data_numerical = data[numerical_features]
```

We will compare the generalization performance of a decision tree and a linear regression. For this purpose, we will create two separate predictive models and evaluate them by 10-fold cross-validation.

Thus,

use `sklearn.linear_model.LinearRegression` and `sklearn.tree.DecisionTreeRegressor` to create the models. Use the default parameters for the linear regression and set `random_state=0` for the decision tree.

Be aware that a linear model requires to scale numerical features. Please use `sklearn.preprocessing.StandardScaler` so that your linear regression model behaves the same way as the quiz author intended ;)

Question 1

(1 point possible)

By comparing the cross-validation test scores for both models fold-to-fold, count the number of times the linear model has a better test score than the decision tree model. Select the range which this number belongs to:

a) [0, 3]: the linear model is substantially worse than the decision tree b) [4, 6]: both models are almost equivalent
c) [7, 10]: the linear model is substantially better than the decision tree

EXPLANATION

solution: c)

The code to get an estimate of the generalization performance of a linear regression is shown below:

```
from sklearn.model_selection import cross_validate  
  
from sklearn.pipeline import make_pipeline  
  
from sklearn.preprocessing import StandardScaler  
  
from sklearn.linear_model import LinearRegression
```

```

linear_regression = make_pipeline(StandardScaler(),
LinearRegression())
cv_results_linear_regression = cross_validate(
    linear_regression, data_numerical, target, cv=10,
return_estimator=True,
    n_jobs=2
)
scores_lr = cv_results_linear_regression["test_score"]
scores_lr

```

We need to preprocess the data with a `StandardScaler` to scale before to train the linear regressor.

When dealing with decision tree, it is unnecessary to scale the data. So we can directly train our decision tree regressor. The code below shows how to evaluate a decision tree regressor on our regression problem.

```

from sklearn.tree import DecisionTreeRegressor

tree = DecisionTreeRegressor(random_state=0)
cv_results_tree = cross_validate(
    tree, data_numerical, target, cv=10, n_jobs=2
)
score_tree = cv_results_tree["test_score"]

```

We can then compare the individual CV scores.

```
print(
```

```
"Linear regression is better than decision tree for "
```

```
f"{sum(scores_lr > score_tree)} CV iterations out of 10 folds."
```

```
)
```

And the linear regression model is better than the tree model.

Linear regression is better than decision tree for 9 CV iterations out of

10 folds.

Hide Answer

You have used 1 of 1 submissions

Question 2

(1/1 point)

Instead of using the default parameters for the decision tree regressor, we can optimize the `max_depth` of the tree. Vary the `max_depth` from 1 level up to 15 levels. Use nested cross-validation to evaluate a grid-search

(`sklearn.model_selection.GridSearchCV`). Set `cv=10` for both the inner and outer cross-validations, then answer the questions below

What is the optimal tree depth for the current problem?

a) The optimal depth is ranging from 3 to 5 b) The optimal depth is ranging from 5 to 8 b) The optimal depth is ranging from 5 to 8 - correct c) The optimal depth is ranging from 8 to 11 d) The optimal depth is ranging from 11 to 15

EXPLANATION

solution: b)

We need to use a `GridSearchCV` and pass a `DecisionTreeRegressor` .

In addition, we should pass an array with the different values

of `max_depth` to evaluate. We give such an example below:

```
import numpy as np

from sklearn.model_selection import GridSearchCV

params = {"max_depth": np.arange(1, 16)}

search = GridSearchCV(tree, params, cv=10)

cv_results_tree_optimal_depth = cross_validate(

    search, data_numerical, target, cv=10, return_estimator=True,

    n_jobs=2,

)
```

We can inspect the trees trained during the cross-validation to find the values

of the `max_depth` parameter.

```
for search_cv in cv_results_tree_optimal_depth["estimator"]:

    print(search_cv.best_params_)
```

Most of the time `max_depth` was in the range 5 to 8.

Hide Answer

You have used 1 of 1 submissions

Question 3

(1/1 point)

Now, we want to evaluate the generalization performance of the decision tree while taking into account the fact that we tune the depth for this specific dataset.

Use the grid-search as an estimator inside a `cross_validate` to automatically tune the `max_depth` parameter on each cross-validation fold.

A tree with tuned depth

a) is always worse than the linear models on all CV folds
b) is often but not always worse than the linear model
c) is often but not always better than the linear model - correct
d) is always better than the linear models on all CV folds

Note: Try to set the `random_state` of the decision tree to different values e.g. `random_state=1` or `random_state=2` and re-run the nested cross-validation to check that your answer is stable enough.

EXPLANATION

solution: b)

```
search = GridSearchCV(tree, params, cv=10)

cv_results_tree_optimal_depth = cross_validate(
    search, data_numerical, target, cv=10, return_estimator=True,
    n_jobs=2,
)

cv_results_tree_optimal_depth["test_score"].mean()
```

We can then check the generalization performance of the model.

```
print(
    "A tree with an optimized depth is better than linear regression for "
    f"{sum(cv_results_tree_optimal_depth['test_score'] > scores_lr)} CV "
    "iterations out of 10 folds."
)
```

A tree with an optimized depth is better than linear regression for 2 CV

iterations out of 10 folds.

Hide Answer

You have used 1 of 1 submissions

Question 4

(1/1 point)

Instead of using only the numerical features you can now use the entire dataset available in the variable `data`.

Create a preprocessor by dealing separately with the numerical and categorical columns. For the sake of simplicity, we assume the following:

- categorical columns can be selected if they have an `object` data type;
- use an `OrdinalEncoder` to encode the categorical columns;
- numerical columns should correspond to the `numerical_features` as defined above.

In addition, set the `max_depth` of the decision tree to `7` (fixed, no need to tune it with a grid-search).

Evaluate this model using `cross_validate` as in the previous questions.

A tree model trained with both numerical and categorical features

a) is most often worse than the tree model using only the numerical features b)
is most often better than the tree model using only the numerical features b) is
most often better than the tree model using only the numerical features -
correct

Note: Try to set the `random_state` of the decision tree to different values e.g. `random_state=1` or `random_state=2` and re-run the (this time single) cross-validation to check that your answer is stable enough.

EXPLANATION

solution: b)

The code to create the predictive model to handle both categorical and numerical columns is the following:

```
from sklearn.compose import make_column_transformer

from sklearn.compose import make_column_selector as selector

from sklearn.preprocessing import OrdinalEncoder

categorical_processor = OrdinalEncoder(

    handle_unknown="use_encoded_value", unknown_value=-1

)

preprocessor = make_column_transformer(

    (categorical_processor, selector(dtype_include=object)),

    ("passthrough", numerical_features)

)

tree = make_pipeline(preprocessor,

    DecisionTreeRegressor(max_depth=7, random_state=0))
```

Then, we can evaluate this tree using cross-validation:

```
cv_results = cross_validate(

    tree, data, target, cv=10, return_estimator=True, n_jobs=2

)

cv_results["test_score"].mean()

import matplotlib.pyplot as plt
```



```
test_score_num = cv_results_tree_optimal_depth["test_score"]
```

```
test_score_all = cv_results["test_score"]
```

```
indices = np.arange(len(test_score_num))
```

```
plt.scatter(
```

```
    indices, test_score_num, color="tab:blue", label="numerical
```

```
features only"
```

```
)
```

```
plt.scatter(
```

```
    indices,
```

```
    test_score_all,
```

```
    color="tab:red",
```

```
    label="all features",
```

```
)
```

```
plt.ylim((0, 1))
```

```
plt.xlabel("Cross-validation iteration")
```

```
plt.ylabel("R2 score")
```

```
_ = plt.legend(loc="lower right")
```

```
print(
```

```
    "A tree model using both numerical and categorical features is better than a "
```

```
    "tree with optimal depth using only numerical features for "
```

```
    f"{sum(cv_results['test_score'] > cv_results_tree_optimal_depth['test_score'])} CV
```

```
"
```

"iterations out of 10 folds."

)

A tree model using both numerical and categorical features is better than a tree with optimal depth using only numerical features for 7 CV iterations out of 10 folds.